

HuGen2071 book

Daniel E. Weeks

August 9, 2022

Table of contents

Preface	5
1 Preparation	6
1.1 Basic programming ideas	6
1.1.1 Introduction to Coding	6
1.2 R	6
1.2.1 PhD Training Workshop: Statistics in R	6
1.3 R and RStudio	6
1.3.1 R for the Rest of Us	6
1.4 GitHub	7
1.5 R Markdown	7
1.6 Unix	7
2 Introduction	8
3 Logistics	9
3.1 GitHub: Set up an account	9
3.2 GitHub Classroom	9
4 R Basics Group Exercise	10
4.1 Set up the data frame a	10
4.2 Exercise 1: recycling	10
4.3 Exercise 2: vector addition	11
4.4 Exercise 3: for loops	12
4.5 Exercise 4: while loops	14
4.6 Exercise 5: repeat loops	15
4.7 Exercise 6: using the rep function	16
4.8 Exercise 7	16
4.9 Exercise 8	17
5 R Character Exercise	20
6 Load Libraries	21
7 Useful RStudio cheatsheet	22

8 Scenario 1	23
9 Discussion Questions	25
9.1 Question 1	25
9.2 Answer 1	25
9.3 Question 2	26
9.4 Answer 2	26
9.5 Question 3	27
9.6 Answer 3	28
10 Scenario 2	29
10.1 Question 4	29
10.2 Answer 4	30
11 R Functions Exercise	32
12 Load Libraries	33
13 Location	34
14 Data set creation code	35
15 Example	36
15.1 Question: How could we construct a list of file names?	37
15.2 Question: Outline a possible algorithm	38
15.3 Question: Construct a more detailed step-by-step algorithm.	38
15.4 Task: Write a <code>read_data_file</code> function.	38
15.5 Question: What does the above code assume?	39
15.6 Question: Extend your function to process all of the files	40
15.7 Bonus question	40
16 R Tidyverse Exercise	42
17 Load Libraries	43
18 Untidy data	44
19 Tidy data	45
20 Gather	46
21 Pivot_longer	47
22 WHO TB data	48

23 Conclusion	51
24 Acknowledgment	52
25 R Recoding Reshaping Exercise	53
26 Load Libraries	54
27 Project 1 Data	55
28 Exercise 1	57
29 Checking for duplicates	59
30 Counting the number of occurrences of the ID	60
31 Count sample_id duplicates	61
32 Checking for duplicates	62
32.1 How to list all duplicates	62
32.2 Sample ID	62
32.3 Subject ID	63
33 Exercise 2	64
33.1 Comment	65
34 Exercise 3	66
34.1 Comment:	66
34.2 xtabs table with labels	67
35 Exercise 4	69
36 Exercise 5: Recoding data	70
37 Recoding data	71
37.1 Comment	71
38 Exercise 6	73
39 Exercise 7	74
40 Summary	75
References	76

Preface

This is a Quarto book.

To learn more about Quarto books visit <https://quarto.org/docs/books>.

1 Preparation

The first part of our HuGen 2071 course aims to teach you R in the context of applied data wrangling in a genetic context. In our experience, if you have never programmed before, it moves kind of fast. As such, it would be useful to review these sources below.

1.1 Basic programming ideas

1.1.1 Introduction to Coding

This web page and two short videos discusses how computer programming is very similar to writing a recipe - you have to break a complex project down into precise smaller individual steps.

<https://subjectguides.york.ac.uk/coding/introduction>

1.2 R

1.2.1 PhD Training Workshop: Statistics in R

This online book has a nice introduction to the concepts of programming, RStudio, and R

https://bookdown.org/animestina/R_Manchester/

See Chapters 1, 2, and 3

1.3 R and RStudio

1.3.1 R for the Rest of Us

Acquaint or refresh yourself with R and RStudio — including installing them on your computer with this “R for the Rest of Us course” (24 min of videos + exercises):

<https://rfortherestofus.com/courses/getting-started/>

Slides: <https://rfortherestofus.github.io/getting-started/slides/slides.html>

1.4 GitHub

To introduce yourself to GitHub:

<https://guides.github.com/introduction/git-handbook/>

<https://guides.github.com/activities/hello-world/>

1.5 R Markdown

To introduce yourself or refresh yourself on R Markdown:

<https://rmarkdown.rstudio.com/> (click on Get Started)

1.6 Unix

And finally, to introduce yourself or refresh yourself with Unix (well, Linux in this case, but close enough), try Lessons 1–11 here:

<https://www.webminal.org/>

2 Introduction

This is a book created from markdown and executable code using Quarto within RStudio.

Book source code: <https://github.com/DanielEWeeks/HuGen2071>

Created by Daniel E. Weeks

Website: <https://www.publichealth.pitt.edu/home/directory/daniel-e-weeks>

3 Logistics

3.1 GitHub: Set up an account

Please go to <https://github.com> and set up a GitHub account.

Choose your GitHub user name carefully, as you may end up using it later in a professional context.

3.2 GitHub Classroom

As GitHub Classroom will be used to distribute course materials and to submit assignments, it would be best if you get git working on your own computer. The easiest way to do this is to install RStudio, R, and git on your computer.

Please follow the detailed instructions in <https://github.com/jfikel/github-classroom-for-students>

In particular, see Step 5 re generating an ssh key so you don't need to login every time.

4 R Basics Group Exercise

4.1 Set up the data frame a

```
a <- data.frame(n = 1:4)
dim(a)
```

```
[1] 4 1
```

```
a
```

```
  n
1 1
2 2
3 3
4 4
```

4.2 Exercise 1: recycling

This exercise should help answer this question: ‘In what type of situations would “recycling” be useful?’

Use recycling to insert into the data frame **a** a column named **rowNum1** that contains a 1 in even rows and a 2 in odd rows.

Tip

The R command

```
a$rowNum1 <- NA
```

would insert a new row into the data frame **a** full of NA values.

💡 Expand to see the answer

```
a$rowNum1 <- c(1,2)
a

  n rowNum1
1 1        1
2 2        2
3 3        1
4 4        2
```

4.3 Exercise 2: vector addition

Use vector addition to construct a vector of length 4 that contains a 1 in even rows and a 2 in odd rows. Then insert this vector into the data frame **a** into a column named **rowNum6**.

💡 Tip

What vector could you add to this vector so the sum is the vector (1, 2, 1, 2)?

```
rep(1, 4)

[1] 1 1 1 1
```

💡 Expand to see the answer

```
r1 <- rep(1, times = 4)
r2 <- rep(c(0,1), times = 2)
r1

[1] 1 1 1 1

r2

[1] 0 1 0 1

r1 + r2

[1] 1 2 1 2
```

```

a$rowNum6 <- r1 + r2
a

  n rowNum1 rowNum6
1 1         1       1
2 2         2       2
3 3         1       1
4 4         2       2

```

4.4 Exercise 3: for loops

Loops allow you to repeat actions on each item from a vector of items.

Here is an example `for` loop, iterating through the values of `i` from 1 to 3:

```

for (i in 1:3) {
  print(paste("i =",i))
}

```

```

[1] "i = 1"
[1] "i = 2"
[1] "i = 3"

```

This does the same thing as this repetitive code:

```

i.vector <- c(1,2,3)
i <- i.vector[1]
print(paste("i =",i))

```

```

[1] "i = 1"

```

```

i <- i.vector[2]
print(paste("i =",i))

```

```

[1] "i = 2"

```

```
i <- i.vector[3]
print(paste("i =",i))
```

```
[1] "i = 3"
```

Use a `for` loop to insert into the data frame `a` a column named `rowNum2` that contains a 1 in even rows and a 2 in odd rows.

Tip

Think about how as `i` increments from 1 to `nrow(a)`, how could we map that sequence (e.g. 1, 2, 3, 4) to the desired sequence of 1, 2, 1, 2.

Expand to see the answer

```
# Set value that we want to iterate 1, 2, 1, 2, ...
j <- 1
# Initialize rowNum2 to all missing values
a$rowNum2 <- NA
# Start the for loop, looping over the number of rows in a
for (i in c(1:nrow(a))) {
  # Assign value j to row i
  a$rowNum2[i] <- j
  # Increment j
  j <- j + 1
  # If j is greater than 2, set it back to 1
  if (j > 2) {
    j <- 1
  }
}
```

```
a
```

	n	rowNum1	rowNum6	rowNum2
1	1	1	1	1
2	2	2	2	2
3	3	1	1	1
4	4	2	2	2


4.5 Exercise 4: while loops

Here's an example while loop:

```
i <- 1
while (i < 4) {
  print(paste("i =",i))
  i <- i + 1
}
```

```
[1] "i = 1"
[1] "i = 2"
[1] "i = 3"
```

Use a while loop to insert into the data frame **a** a column named **rowNum3** that contains a 1 in even rows and a 2 in odd rows.

 Expand to see the answer

```
a$rowNum3 = NA
i <- 1 #set index
while(i <= nrow(a)){ #set conditions for while loop

  if ((i %% 2)) { #if statement for when "i" is odd
    a$rowNum3[i] <- 1
  }
  else #else statement for when "i" is even
    a$rowNum3[i] <- 2

  i <- i + 1 #counter for "i", increments by 1 with each loop iteration
}
a
```

	n	rowNum1	rowNum6	rowNum2	rowNum3
1	1	1	1	1	1
2	2	2	2	2	2
3	3	1	1	1	1
4	4	2	2	2	2


4.6 Exercise 5: repeat loops

Here's an example `repeat` loop:

```
i <- 1
repeat {
  print(paste("i =",i))
  i <- i + 1
  if (i > 3) break
}
```

```
[1] "i = 1"
[1] "i = 2"
[1] "i = 3"
```

Use a `repeat` loop to insert into the data frame `a` a column named `rowNum4` that contains a 1 in even rows and a 2 in odd rows.

 Expand to see the answer

```
a$rowNum4 <- NA
i <- 1 #set index
repeat {

  if ((i %% 2)) { #if statement for when "i" is odd
    a$rowNum4[i] <- 1
  }
  else #else statement for when "i" is even
    a$rowNum4[i] <- 2

  i <- i + 1 #counter for "i", increments by 1 with each loop iteration
  if (i > nrow(a)) {
    break
  }
}
a
```

	n	rowNum1	rowNum6	rowNum2	rowNum3	rowNum4
1	1	1	1	1	1	1
2	2	2	2	2	2	2
3	3	1	1	1	1	1

4	4	2	2	2	2	2
---	---	---	---	---	---	---

4.7 Exercise 6: using the rep function

Use the `rep` command to insert into the data frame `a` a column named `rowNum5` that contains a 1 in even rows and a 2 in odd rows.

💡 Expand to see the answer

```
# This will only work correctly if nrow(a) is even
a$rowNum5 <- rep(c(1,2), nrow(a)/2)
a
```

	n	rowNum1	rowNum6	rowNum2	rowNum3	rowNum4	rowNum5
1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2
3	3	1	1	1	1	1	1
4	4	2	2	2	2	2	2

4.8 Exercise 7

List all even rows of the data frame `a`.

List rows 3 and 4 of the data frame `a`.

💡 Expand to see the answer

```
# All even rows
a[a$rowNum1==2,]

n rowNum1 rowNum6 rowNum2 rowNum3 rowNum4 rowNum5
2 2      2      2      2      2      2      2
4 4      2      2      2      2      2      2

# All odd rows
a[a$rowNum1==1,]

n rowNum1 rowNum6 rowNum2 rowNum3 rowNum4 rowNum5
```


1	1	1	1	1	1	1
3	3	1	1	1	1	1

4.9 Exercise 8

Note

Learning objective: Learn how to alter the options of an R command to achieve your goals.

This exercise should help answer this question: “When reading a file, will missing data be automatically represented as NA values, or does that need to be coded/manually curated?”

The tab-delimited file in `testdata.txt` contains the following data:

1	1	1
2	2	2
3	NA	99
4	4	4

Your collaborator who gave you these data informed you that in this file 99 stands for a missing value, as does NA.

However if we use the `read.table` command with its default options to read this in, we fail to accomplish the desired task, as 99 is not reading as a missing value:

```
infile <- "data/testdata.txt"
# Adjust the read.table options to read the file correctly as desired.
b <- read.table(infile)
b
```

	V1	V2	V3
1	1	1	1
2	2	2	2
3	3	NA	99
4	4	4	4

```
str(b)
```

```
'data.frame':  4 obs. of  3 variables:
 $ V1: int  1 2 3 4
 $ V2: int  1 2 NA 4
 $ V3: int  1 2 99 4
```

Use the `read.table` command to read this file in while automatically setting both the 'NA' and the 99 to NA. This can be done by adjusting the various options of the `read.table` command.

Tip

Read the help page for the `read.table` command

Expand to see the answer

To read this in properly, we have to let 'read.table' know that there is no header and that which values should be mapped to the missing NA value:

```
b <- read.table(infile, header = FALSE, na.strings = c("NA","99"))
b
```

```
  V1 V2 V3
1  1  1  1
2  2  2  2
3  3 NA NA
4  4  4  4
```

```
str(b)
```

```
'data.frame':  4 obs. of  3 variables:
 $ V1: int  1 2 3 4
 $ V2: int  1 2 NA 4
 $ V3: int  1 2 NA 4
```

```
summary(b)
```

	V1	V2	V3
Min.	:1.00	Min. :1.000	Min. :1.000
1st Qu.:	1.75	1st Qu.:1.500	1st Qu.:1.500
Median	:2.50	Median :2.000	Median :2.000
Mean	:2.50	Mean :2.333	Mean :2.333

3rd Qu.:3.25	3rd Qu.:3.000	3rd Qu.:3.000
Max. :4.00	Max. :4.000	Max. :4.000
	NA's :1	NA's :1

5 R Character Exercise

6 Load Libraries

```
library(tidyverse)
# library(tidylog)
library(knitr)
```

7 Useful RStudio cheatsheet

See the “String manipulation with stringr cheatsheet” at
<https://www.rstudio.com/resources/cheatsheets/>

8 Scenario 1

You are working with three different sets of collaborators: 1) the clinical group that did the field work and generated the anthropometric measurements; 2) the medical laboratory that measured blood pressure in a controlled environment; and 3) the molecular laboratory that generated the genotypes.

```
clin <- read.table(file = "data/clinical_data.txt", header=TRUE)
kable(clin)
```

ID	height
1	152
104	172
2112	180
2543	163

```
lab <- read.table(file = "data/lab_data.txt", header = TRUE)
kable(lab)
```

ID	SBP
SG0001	120
SG0104	111
SG2112	125
SG2543	119

```
geno <- read.table(file = "data/genotype_data.txt", header = TRUE)
kable(geno)
```

Sample	rs1212
TaqMan-SG0001-190601	G/C
TaqMan-SG0104-190602	G/G
TaqMan-SG2112-190603	C/C

Sample	rs1212
TaqMan-Sg2543-190603	C/G

9 Discussion Questions

9.1 Question 1

The clinical group, which measured height, used integer IDs, but the medical group, which measured the blood pressure, decided to prefix the integer IDs with the string ‘SG’ (so as to distinguish them from other studies that were also using integer IDs). So ID ‘1’ was mapped to ID ‘SG0001’.

Discuss how, using R commands, you would reformat the integer IDs to be in the format “SGXXXX”. Write down your ideas in the next section, and, if you have time, try them out within an R chunk.

Table 9.1: The `clin` data frame

ID	height
1	152
104	172
2112	180
2543	163

Hint: Use the `formatC` function.

9.2 Answer 1

💡 Expand to see solution

```
clin$SUBJECT_ID <- paste0("SG", formatC(clin$ID, width = 4, flag = "0000"))
kable(clin)
```

ID	height	SUBJECT_ID
1	152	SG0001

104	172	SG0104
2112	180	SG2112
2543	163	SG2543

9.3 Question 2

Discuss how, using R commands, you would reformat the “SGXXXX” IDs to be integer IDs. Write down your ideas in the next section, and, if you have time, try them out within an R chunk.

Table 9.3: The `lab` data frame

ID	SBP
SG0001	120
SG0104	111
SG2112	125
SG2543	119

Hint: Use either the `gsub` command or the `str_replace_all` command.

9.4 Answer 2

💡 Expand to see solution

```
lab$ID2 <- as.numeric(gsub("SG","",lab$ID))
kable(lab)
```

ID	SBP	ID2
SG0001	120	1
SG0104	111	104
SG2112	125	2112
SG2543	119	2543

```
lab$ID2 <- NA
lab$ID2 <- str_replace_all(lab$ID, pattern = "SG", replacement = "") %>% as.numeric()
kable(lab)
```

ID	SBP	ID2
SG0001	120	1
SG0104	111	104
SG2112	125	2112
SG2543	119	2543

9.5 Question 3

The genotype group used IDs in the style “TaqMan-SG0001-190601”, where the first string is “TaqMan” and the ending string is the date of the genotyping experiment.

Discuss how, using R commands, you would extract an “SGXXXX” style ID from the “TaqMan-SG0001-190601” style IDs. Write down your ideas in the next section, and, if you have time, try them out within an R chunk.

Note that one of the IDs has a lower case ‘g’ in it - how would you correct this, using R commands?

Table 9.6: The **geno** data frame

Sample	rs1212
TaqMan-SG0001-190601	G/C
TaqMan-SG0104-190602	G/G
TaqMan-SG2112-190603	C/C
TaqMan-Sg2543-190603	C/G

Hint: Use either the `str_split_fixed` function or the `separate` function.

9.6 Answer 3

💡 Expand to see solution

```
a <- str_split_fixed(geno$Sample, pattern = "-", n=3)
a
```

```
      [,1]      [,2]      [,3]
[1,] "TaqMan" "SG0001" "190601"
[2,] "TaqMan" "SG0104" "190602"
[3,] "TaqMan" "SG2112" "190603"
[4,] "TaqMan" "Sg2543" "190603"
```

```
geno$ID <- toupper(a[,2])
kable(geno)
```

Sample	rs1212	ID
TaqMan-SG0001-190601	G/C	SG0001
TaqMan-SG0104-190602	G/G	SG0104
TaqMan-SG2112-190603	C/C	SG2112
TaqMan-Sg2543-190603	C/G	SG2543

The `separate` function from the `tidyr` package is also useful:

```
geno %>% separate(Sample, into=c("Tech","ID","Suffix"), sep="-")
```

```
      Tech      ID Suffix rs1212
1 TaqMan SG0001 190601    G/C
2 TaqMan SG0104 190602    G/G
3 TaqMan SG2112 190603    C/C
4 TaqMan Sg2543 190603    C/G
```

10 Scenario 2

A replication sample has been measured, and that is using IDs in the style “RP5XXX”.

```
joint <- read.table(file = "data/joint_data.txt", header = TRUE)
kable(joint)
```

ID	SBP
SG0001	120
SG0104	111
SG2112	125
SG2543	119
RP5002	121
RP5012	118
RP5113	112
RP5213	142

10.1 Question 4

Discuss how you would use R commands to split the ‘joint’ data frame into an ‘SG’ and ‘RP’ specific piece? Write down your ideas in the next section, and, if you have time, try them out within an R chunk.

Table 10.2: The joint data frame

ID	SBP
SG0001	120
SG0104	111
SG2112	125
SG2543	119
RP5002	121
RP5012	118
RP5113	112
RP5213	142

10.2 Answer 4

💡 Expand to see solution

```
grep(pattern = "SG", joint$ID)
```

```
[1] 1 2 3 4
```

```
grep(pattern = "RP", joint$ID)
```

```
[1] 5 6 7 8
```

```
joint.SG <- joint[grep(pattern = "SG", joint$ID), ]  
joint.RP <- joint[grep(pattern = "RP", joint$ID), ]  
kable(joint.SG)
```

ID	SBP
SG0001	120
SG0104	111
SG2112	125
SG2543	119

```
kable(joint.RP)
```

	ID	SBP
5	RP5002	121
6	RP5012	118
7	RP5113	112
8	RP5213	142

```
# Reset row names  
rownames(joint.RP) <- NULL  
kable(joint.RP)
```

ID	SBP
RP5002	121
RP5012	118
RP5113	112
RP5213	142

11 R Functions Exercise

12 Load Libraries

```
library(tidyverse)
# library(tidylog)
```

13 Location

This file `exercise1_solution.Qmd` is in the “HuGen2071__book” sub-folder of the “hugen2071” folder of our Lectures repository.

```
paste0(basename(dirname(getwd())),"/",basename(getwd()))
```

```
[1] "hugen2071/HuGen2071_book"
```

14 Data set creation code

```
i <- 6
for (i in 1:10) {
  fl <- data.frame(name=rep(paste0("name",i),26))
  b <- data.frame(name = rep(NA, 26))
  b$name <- paste0(fl$name,"_",letters)
  b$trait <- rnorm(26)
  write_tsv(b,paste0("data/dataset",i,".txt"))
}
```

15 Example

Here we have been sent three data sets in the files that contain the trait quantitative values for each person in the data set:

“dataset1.txt” “dataset2.txt” “dataset3.txt”

And we’ve been asked to make a table that gives, for each dataset, the sample size (N), the mean of the trait, the median, and the variance.

We could do this by reading in each data set, one by one, as follows:

```
results <- data.frame(dataset=rep(NA,3),N=NA, mean=NA, median=NA, var=NA)
fl1 <- read.table("data/dataset1.txt",sep="\t",header=TRUE)
results$dataset[1] <- "dataset1"
results$N <- nrow(fl1)
results$mean[1] <- mean(fl1$trait)
results$median[1] <- median(fl1$trait)
results$var[1] <- var(fl1$trait)
results
```

	dataset	N	mean	median	var
1	dataset1	26	0.09762111	0.2198957	0.5974116
2	<NA>	26	NA	NA	NA
3	<NA>	26	NA	NA	NA

```
fl2 <- read.table("data/dataset2.txt",sep="\t",header=TRUE)
results$dataset[2] <- "dataset2"
results$N <- nrow(fl2)
results$mean[2] <- mean(fl2$trait)
results$median[2] <- median(fl2$trait)
results$var[2] <- var(fl2$trait)
results
```

	dataset	N	mean	median	var
1	dataset1	26	0.09762111	0.2198957	0.5974116

```
2 dataset2 26 0.43486401 0.3558736 1.0936651
3    <NA> 26          NA          NA          NA
```

```
fl3 <- read.table("data/dataset3.txt",sep="\t",header=TRUE)
results$dataset[3] <- "dataset3"
results$N <- nrow(fl3)
results$mean[3] <- mean(fl3$trait)
results$median[3] <- median(fl3$trait)
results$var[3] <- var(fl3$trait)
results
```


```
   dataset  N      mean    median      var
1 dataset1 26 0.09762111 0.2198957 0.5974116
2 dataset2 26 0.43486401 0.3558736 1.0936651
3 dataset3 26 0.07508335 0.0445614 0.7950574
```

Your colleague initially sent you the three data sets above, but now your colleague has sent you three more data sets and asked you to update the ‘results’ table.

As you can see, the code above is very repetitive. So let’s automate this by writing a function that loops through a list of data set files named “dataset1.txt”, “dataset2.txt”, “dataset3.txt”, etc., building up the results table as above.

15.1 Question: How could we construct a list of file names?

How could we construct a list of file names?

 Expand to see solution

Hint: the `list.files` command provides a handy way to get a list of the input files:

```
fls <- list.files(path="data",pattern="dataset*")
fls
```

```
[1] "dataset1.txt" "dataset2.txt" "dataset3.txt" "dataset4.txt" "dataset5.txt"
[6] "dataset6.txt"
```

15.2 Question: Outline a possible algorithm

Outline a possible algorithm that loops through a list of input data set files named “dataset1.txt”, “dataset2.txt”, “dataset3.txt”, etc., building up the results table as above.

💡 Expand to see solution

- Read in the input file names into a list
- Set up an empty results table
- For each file in our file name list
 - Read the file
 - Compute the statistics
 - Insert the information into the results table
 - Return the filled-in results table

15.3 Question: Construct a more detailed step-by-step algorithm.

Construct a more detailed step-by-step algorithm.

💡 Expand to see solution

- Input the path to the folder containing the data files
- Read in the input file names into a list `fls`
- Count the number of input files `N`
- Set up an empty results table with `N` rows
- For each file in our file name list `fls`
 - Read the file
 - Compute the statistics
 - Insert the information into the correct row of the results table
- Return the filled-in results table

15.4 Task: Write a `read_data_file` function.

Write a `read_data_file` function to accomplish the required steps for a single input data file.

1. Make the number in the data file name an argument.

💡 Expand to see solution

Here we make the number in the data file name an argument

```
results <- data.frame(dataset=rep(NA,6),N=NA, mean=NA, median=NA, var=NA)
read_data_file <- function(n=1, results) {
  fl1 <- read.table(paste0("data/dataset",n,".txt"),sep="\t",header=TRUE)
  results$dataset[n] <- paste0("dataset",n,".txt")
  results$N <- nrow(fl1)
  results$mean[n] <- mean(fl1$trait)
  results$median[n] <- median(fl1$trait)
  results$var[n] <- var(fl1$trait)
  invisible(results)
}
```

2. Make the path to the input file an argument to your `read_data_file` function.

💡 Expand to see solution

Here we make the path to the input file an argument.

```
read_data_file_v2 <- function(flnm, results) {
  fl1 <- read.table(paste0("data/",flnm),sep="\t",header=TRUE)
  results$dataset[n] <- flnm
  results$N <- nrow(fl1)
  results$mean[n] <- mean(fl1$trait)
  results$median[n] <- median(fl1$trait)
  results$var[n] <- var(fl1$trait)
  invisible(results)
}
```

15.5 Question: What does the above code assume?

What does the above code assume?

💡 Expand to see solution

Assumes a file naming style of 'dataset*.txt' where the asterisk represents 1, 2, 3, ...
Assumes the files are in the "data" folder.

15.6 Question: Extend your function to process all of the files

The above function `read_data_file` processes one file at a time. How would you write a function to loop this over to process all of our files?

💡 Expand to see solution

```
fls <- list.files(path="data",pattern="dataset*")

loop_over_dataset <- function(fls) {
  # Input: the list of file names
  # Output: the 'results table'
  # Count the number of data set file names in fls
  n_datasets <- length(fls)
  # Set up a results dataframe with n_datasets rows
  results <- data.frame(dataset=rep(NA,n_datasets),N=NA, mean=NA, median=NA, var=NA)
  for (n in 1:n_datasets) {
    results <- read_data_file(n=n, results=results)
  }
  return(results)
}

loop_over_dataset(fls = fls)
```

	dataset	N	mean	median	var
1	dataset1.txt	26	0.09762111	0.21989574	0.5974116
2	dataset2.txt	26	0.43486401	0.35587359	1.0936651
3	dataset3.txt	26	0.07508335	0.04456140	0.7950574
4	dataset4.txt	26	0.06259720	0.04813915	0.9186042
5	dataset5.txt	26	-0.09288522	-0.19155759	0.9978161
6	dataset6.txt	26	-0.20266667	-0.23845426	1.5605823

15.7 Bonus question

Can you find a subtle mistake in the `read_data_file` function?

```
results <- data.frame(dataset=rep(NA,6),N=NA, mean=NA, median=NA, var=NA)
read_data_file <- function(n=1, results) {
  fl1 <- read.table(paste0("data/dataset",n,".txt"),sep="\t",header=TRUE)
```



```

results$dataset[n] <- paste0("dataset",n,".txt")
results$N <- nrow(fl1)
results$mean[n] <- mean(fl1$trait)
results$median[n] <- median(fl1$trait)
results$var[n] <- var(fl1$trait)
invisible(results)
}

```

💡 Expand to see solution

If N varies across the data sets, then this line will not do the right thing:

```

results$N <- nrow(fl1)

results <- data.frame(dataset=rep(NA,6),N=NA, mean=NA, median=NA, var=NA)
read_data_file <- function(n=1, results) {
  fl1 <- read.table(paste0("data/dataset",n,".txt"),sep="\t",header=TRUE)
  results$dataset[n] <- paste0("dataset",n,".txt")
  results$N[n] <- nrow(fl1)
  results$mean[n] <- mean(fl1$trait)
  results$median[n] <- median(fl1$trait)
  results$var[n] <- var(fl1$trait)
  invisible(results)
}

```

16 R Tidyverse Exercise

17 Load Libraries

Load the tidyverse packages

```
library(tidyverse)
# library(tidylog)
```

18 Untidy data

Let's use the World Health Organization TB data set from the `tidyr` package

```
who <- tidyr::who
dim(who)
```

```
[1] 7240    60
```

```
head(who[,1:6] %>% filter(!is.na(new_sp_m014)))
```

```
# A tibble: 6 x 6
  country    iso2 iso3  year new_sp_m014 new_sp_m1524
  <chr>      <chr> <chr> <int>      <int>      <int>
1 Afghanistan AF    AFG  1997         0         10
2 Afghanistan AF    AFG  1998        30        129
3 Afghanistan AF    AFG  1999         8         55
4 Afghanistan AF    AFG  2000        52        228
5 Afghanistan AF    AFG  2001       129        379
6 Afghanistan AF    AFG  2002        90        476
```

See the help page for `who` for more information about this data set.

In particular, note this description:

“The data uses the original codes given by the World Health Organization. The column names for columns five through 60 are made by combining `new_` to a code for method of diagnosis (rel = relapse, sn = negative pulmonary smear, sp = positive pulmonary smear, ep = extrapulmonary) to a code for gender (f = female, m = male) to a code for age group (014 = 0-14 yrs of age, 1524 = 15-24 years of age, 2534 = 25 to 34 years of age, 3544 = 35 to 44 years of age, 4554 = 45 to 54 years of age, 5564 = 55 to 64 years of age, 65 = 65 years of age or older).”

So `new_sp_m014` represents the counts of new TB cases detected by a positive pulmonary smear in males in the 0-14 age group.

19 Tidy data

Tidy data: Have each variable in a column.

Question: Are these data tidy?

💡 Expand to see solution

No these data are not tidy because aspects of the data that should be variables are encoded in the name of the variables.

These aspects are

1. test type.
2. sex of the subjects.
3. age range of the subjects.

Question: How would we make these data tidy?

Consider this portion of the data:

```
head(who[,1:5] %>% filter(!is.na(new_sp_m014) & new_sp_m014>0), 1)
```

```
# A tibble: 1 x 5
  country    iso2 iso3  year new_sp_m014
  <chr>      <chr> <chr> <int>      <int>
1 Afghanistan AF    AFG   1998         30
```

💡 Expand to see solution

We would replace the `new_sp_m014` with the following four columns:

```
type  sex  age  n
sp    m   014 30
```

This would place each variable in its own column.

20 Gather

```
stocks <- tibble(  
  time = as.Date('2009-01-01') + 0:9,  
  X = rnorm(10, 0, 1),  
  Y = rnorm(10, 0, 2),  
  Z = rnorm(10, 0, 4)  
)  
  
head(stocks)
```

```
# A tibble: 6 x 4  
  time           X         Y         Z  
  <date>       <dbl>   <dbl> <dbl>  
1 2009-01-01 -0.0999 -0.643 -3.73  
2 2009-01-02 -0.314  -0.0374  4.30  
3 2009-01-03  0.900   1.67   -8.03  
4 2009-01-04 -0.671   0.506  -1.69  
5 2009-01-05  0.483   0.490   3.79  
6 2009-01-06  0.907   5.59   9.98
```

```
stocks %>% gather("stock", "price", -time) %>% head()
```

```
# A tibble: 6 x 3  
  time      stock  price  
  <date>   <chr>   <dbl>  
1 2009-01-01 X      -0.0999  
2 2009-01-02 X      -0.314  
3 2009-01-03 X       0.900  
4 2009-01-04 X     -0.671  
5 2009-01-05 X       0.483  
6 2009-01-06 X       0.907
```

21 Pivot_longer

```
stocks %>% pivot_longer(c(X,Y,Z), names_to= "stock", values_to = "price") %>%  
  head()
```

```
# A tibble: 6 x 3  
  time      stock  price  
  <date>    <chr>   <dbl>  
1 2009-01-01 X      -0.0999  
2 2009-01-01 Y      -0.643  
3 2009-01-01 Z      -3.73  
4 2009-01-02 X      -0.314  
5 2009-01-02 Y      -0.0374  
6 2009-01-02 Z       4.30
```

22 WHO TB data

Question: How would we convert this to tidy form?

```
head(who[,1:6] %>% filter(!is.na(new_sp_m014)))
```

```
# A tibble: 6 x 6
  country    iso2 iso3  year new_sp_m014 new_sp_m1524
  <chr>      <chr> <chr> <int>      <int>      <int>
1 Afghanistan AF    AFG   1997         0         10
2 Afghanistan AF    AFG   1998        30        129
3 Afghanistan AF    AFG   1999         8         55
4 Afghanistan AF    AFG   2000        52        228
5 Afghanistan AF    AFG   2001       129        379
6 Afghanistan AF    AFG   2002        90        476
```

💡 Expand to see solution

```
who.long <- who %>% pivot_longer(starts_with("new"), names_to = "demo", values_to = "n")
head(who.long)
```

```
# A tibble: 6 x 6
  country    iso2 iso3  year demo          n
  <chr>      <chr> <chr> <int> <chr>      <int>
1 Afghanistan AF    AFG   1997 new_sp_m014    0
2 Afghanistan AF    AFG   1997 new_sp_m1524   10
3 Afghanistan AF    AFG   1997 new_sp_m2534    6
4 Afghanistan AF    AFG   1997 new_sp_m3544    3
5 Afghanistan AF    AFG   1997 new_sp_m4554    5
6 Afghanistan AF    AFG   1997 new_sp_m5564    2
```

Question: How would we split demo into variables?

```
head(who.long)
```



```
# A tibble: 6 x 6
  country    iso2 iso3  year demo          n
  <chr>      <chr> <chr> <int> <chr>      <int>
1 Afghanistan AF    AFG   1997 new_sp_m014    0
2 Afghanistan AF    AFG   1997 new_sp_m1524   10
3 Afghanistan AF    AFG   1997 new_sp_m2534    6
4 Afghanistan AF    AFG   1997 new_sp_m3544    3
5 Afghanistan AF    AFG   1997 new_sp_m4554    5
6 Afghanistan AF    AFG   1997 new_sp_m5564    2
```

Look at the variable naming scheme:

```
names(who) %>% grep("m014",., value=TRUE)
```

```
[1] "new_sp_m014" "new_sn_m014" "new_ep_m014" "newrel_m014"
```

Question: How should we adjust the `demo` strings so as to be able to easily split all of them into the desired variables?

 Expand to see solution

```
who.long <- who.long %>%
  mutate(demo = str_replace(demo, "newrel", "new_rel"))
grep("m014",who.long$demo, value=TRUE) %>% unique()
```

```
[1] "new_sp_m014" "new_sn_m014" "new_ep_m014" "new_rel_m014"
```

Question: After adjusting the `demo` strings, how would we then separate them into the desired variables?

 Expand to see solution

```
who.long <- who.long %>%
  separate(demo, into = c("new", "type", "sexagerange"), sep="_") %>%
  separate(sexagerange, into=c("sex","age_range"), sep=1) %>%
  select(-new)
head(who.long)
```

```
# A tibble: 6 x 8
  country    iso2 iso3  year type  sex  age_range    n
  <chr>      <chr> <chr> <int> <chr> <chr> <chr>      <int>
1 Afghanistan AF    AFG   1997 new  sp  new_sp_m014    0
2 Afghanistan AF    AFG   1997 new  sn  new_sn_m014   10
3 Afghanistan AF    AFG   1997 new  ep  new_ep_m014    6
4 Afghanistan AF    AFG   1997 new  rel new_rel_m014    3
5 Afghanistan AF    AFG   1997 new  rel new_rel_m014    5
6 Afghanistan AF    AFG   1997 new  rel new_rel_m014    2
```

	<chr>	<chr>	<chr>	<int>	<chr>	<chr>	<chr>	<int>
1	Afghanistan	AF	AFG	1997	sp	m	014	0
2	Afghanistan	AF	AFG	1997	sp	m	1524	10
3	Afghanistan	AF	AFG	1997	sp	m	2534	6
4	Afghanistan	AF	AFG	1997	sp	m	3544	3
5	Afghanistan	AF	AFG	1997	sp	m	4554	5
6	Afghanistan	AF	AFG	1997	sp	m	5564	2

23 Conclusion

Now our untidy data are tidy.

```
head(who.long)
```

```
# A tibble: 6 x 8
  country    iso2 iso3  year type sex age_range     n
  <chr>      <chr> <chr> <int> <chr> <chr> <chr>   <int>
1 Afghanistan AF    AFG  1997 sp    m    014         0
2 Afghanistan AF    AFG  1997 sp    m   1524        10
3 Afghanistan AF    AFG  1997 sp    m   2534         6
4 Afghanistan AF    AFG  1997 sp    m   3544         3
5 Afghanistan AF    AFG  1997 sp    m   4554         5
6 Afghanistan AF    AFG  1997 sp    m   5564         2
```

24 Acknowledgment

This exercise was modeled, in part, on this exercise:

https://people.duke.edu/~ccc14/cfar-data-workshop-2018/CFAR_R_Workshop_2018_Exercisees.html

25 R Recoding Reshaping Exercise

26 Load Libraries

```
library(tidyverse)
# library(tidylog)
```

27 Project 1 Data

In the `ds` data frame we have the synthetic yet realistic data we will be using in Project 1.

In the `dd` data frame we have the corresponding data dictionary.

```
load("data/exercise.RData", verbose = TRUE)
```

Loading objects:

```
ds
dd
DictPer
```

```
dim(ds)
```

```
[1] 191 24
```

```
names(ds)
```

```
[1] "sample_id"           "Sample_trimester"
[3] "Gestationalage_sample" "subject_id"
[5] "strata"              "race"
[7] "maternal_age_delivery" "case_control_status"
[9] "prepregnancy_weight" "height"
[11] "prepregnancy_BMI"    "gravidity"
[13] "parity"              "gestationalage_delivery"
[15] "average_SBP_lt20weeks" "average_DBP_lt20weeks"
[17] "average_SBP_labor"   "average_DBP_labor"
[19] "smoke_lifetime"      "baby_birthweight"
[21] "baby_sex"            "baby_birthweight_centile"
[23] "baby_SGA"            "placental_pathology"
```

```
dim(dd)
```

```
[1] 27 5
```

```
names(dd)
```

```
[1] "Original.Variable.Name" "R21.Variable.Name" "Description"
[4] "Variable.Units"       "Variable.Coding"
```



28 Exercise 1

Skill: Checking for duplicated IDs

```
ds %>%  
  select(subject_id, sample_id, height) %>%  
  head(n = 10)
```

	subject_id	sample_id	height
1	SUBJ48	SAMP149	64.6
2	SUBJ46	SAMP037	65.7
3	SUBJ28	SAMP120	63.3
4	SUBJ26	SAMP187	61.1
5	SUBJ49	SAMP082	67.6
6	SUBJ48	SAMP149	64.6
7	SUBJ19	SAMP074	66.1
8	SUBJ07	SAMP063	64.4
9	SUBJ28	SAMP053	63.3
10	SUBJ43	SAMP085	65.7


Check if there are any duplicated `sample_id`'s using the `duplicated` command.

 Expand to see solution

```
sum(duplicated(ds$sample_id))
```

```
[1] 72
```

Construct a table of the number of times each `sample_id` is duplicated:

 Expand to see solution

```
table(table(ds$sample_id))
```

```

 1  2  3  4  5
67 35 13  2  1

# But?
sum(duplicated(ds$sample_id))

[1] 72

35 + 13 * 2 + 2 * 3 + 1 * 4

[1] 71

sum(duplicated(ds$sample_id, incomparables = NA))

[1] 71

table(table(ds$sample_id, useNA = "always"))


 1  2  3  4  5
67 36 13  2  1

36 + 13 * 2 + 2 * 3 + 1 * 4

[1] 72

```

Check if there are any duplicated `subject_ids`

 Expand to see solution

```

sum(duplicated(ds$subject_id))

[1] 137

```

29 Checking for duplicates

How do we return every row that contains a duplicate?

```
f <- data.frame(ID = c(1, 1, 2), c2 = c(1, 2, 3))  
f
```

	ID	c2
1	1	1
2	1	2
3	2	3

```
f[duplicated(f$ID), ]
```

	ID	c2
2	1	2

30 Counting the number of occurrences of the ID

```
f %>%  
  group_by(ID) %>%  
  summarise(n = n())
```

```
# A tibble: 2 x 2  
  ID     n  
<dbl> <int>  
1     1     2  
2     2     1
```

31 Count sample_id duplicates

Using Tidyverse commands, count how many times each `sample_id` occurs in the `ds` data frame, reporting the counts in descending order, from highest to lowest.

💡 Expand to see solution

```
ds %>%  
  group_by(sample_id) %>%  
  summarise(n = n()) %>%  
  filter(n > 1) %>%  
  arrange(desc(n)) %>%  
  head()
```

```
# A tibble: 6 x 2  
  sample_id      n  
  <chr>      <int>  
1 SAMP100        5  
2 SAMP125        4  
3 SAMP155        4  
4 SAMP017        3  
5 SAMP048        3  
6 SAMP058        3
```

```
ds %>%  
  group_by(sample_id) %>%  
  summarise(n = n()) %>%  
  filter(n > 1) %>%  
  arrange(desc(n)) %>%  
  pull(n) %>%  
  table()
```

```
.  
 2  3  4  5  
36 13  2  1
```

32 Checking for duplicates


Here we list all of the rows containing a duplicated 'ID' value using functions from the 'tidyverse' package:

```
f %>%  
  group_by(ID) %>%  
  filter(n() > 1)
```

```
# A tibble: 2 x 2  
# Groups:   ID [1]  
   ID     c2  
<dbl> <dbl>  
1     1     1  
2     1     2
```

32.1 How to list all duplicates

Use Tidyverse commands to list all duplicates for `sample_id` and for `subject_id`. Sort the results by the ID.

 Expand to see solution

32.2 Sample ID

```
ds %>%  
  group_by(sample_id) %>%  
  filter(n() > 1) %>%  
  select(sample_id, subject_id, Sample_trimester, Gestationalage_sample) %>%  
  arrange(sample_id, Sample_trimester, Gestationalage_sample) %>%  
  head()
```

```
# A tibble: 6 x 4  
# Groups:   sample_id [3]
```

	sample_id	subject_id	Sample_trimester	Gestationalage_sample
	<chr>	<chr>	<dbl>	<dbl>
1	SAMP002	SUBJ20	2	19.3
2	SAMP002	SUBJ20	2	19.7
3	SAMP003	SUBJ12	1	8.25
4	SAMP003	SUBJ12	1	8.35
5	SAMP004	SUBJ35	2	20.4
6	SAMP004	SUBJ35	2	20.9

32.3 Subject ID

```
ds %>%
  group_by(subject_id) %>%
  filter(n() > 1) %>%
  select(subject_id, sample_id, Sample_trimester, Gestationalage_sample) %>%
  arrange(subject_id, sample_id, Sample_trimester, Gestationalage_sample) %>%
  head(10)
```

A tibble: 10 x 4


Groups: subject_id [2]

	subject_id	sample_id	Sample_trimester	Gestationalage_sample
	<chr>	<chr>	<dbl>	<dbl>
1	SUBJ01	SAMP011	1	9.00
2	SUBJ01	SAMP034	3	39.8
3	SUBJ01	SAMP034	3	42.1
4	SUBJ01	SAMP103	2	19.9
5	SUBJ01	SAMP103	2	20.0
6	SUBJ01	SAMP155	3	40.0
7	SUBJ01	SAMP155	3	40.5
8	SUBJ01	SAMP155	3	40.7
9	SUBJ01	SAMP155	3	41.6
10	SUBJ02	SAMP113	3	38.6

33 Exercise 2

Skill: Reshaping data

Select only three columns “sample_id”, “Sample_trimester”, “Gestationalage_sample”, and then reshape from ‘long’ format to ‘wide’ format using `pivot_wider`, taking time as the “Sample_trimester”.

 Expand to see solution

```
b <- ds %>%
  select(sample_id, Sample_trimester, Gestationalage_sample)

b2 <- b %>%
  pivot_wider(id_cols = sample_id, names_from = Sample_trimester, values_from = Gestationalage_sample)
```

Warning: Values from `Gestationalage_sample` are not uniquely identified; output will contain data from many groups for some identifiers.

- * Use `values_fn = list` to suppress this warning.
- * Use `values_fn = {summary_fun}` to summarise duplicates.
- * Use the following dplyr code to identify duplicates.

```
{data} %>%
  dplyr::group_by(sample_id, Sample_trimester) %>%
  dplyr::summarise(n = dplyr::n(), .groups = "drop") %>%
  dplyr::filter(n > 1L)
```

```
head(b2)
```

```
# A tibble: 6 x 5
  sample_id `1`      `3`      `2`      `NA`
  <chr>      <list>    <list>    <list>    <list>
1 SAMP149    <dbl [3]> <NULL>    <NULL>    <NULL>
2 SAMP037    <dbl [2]> <NULL>    <NULL>    <NULL>
3 SAMP120    <dbl [3]> <NULL>    <NULL>    <NULL>
4 SAMP187    <NULL>    <dbl [1]> <NULL>    <NULL>
5 SAMP082    <NULL>    <NULL>    <dbl [1]> <NULL>
6 SAMP074    <NULL>    <NULL>    <dbl [3]> <NULL>
```


33.1 Comment

View `b2` via the `View(b2)` command in RStudio - it nicely put all the different gestational age observations into one list for each `sample_id` x `Sample_trimester` combination.

34 Exercise 3

Skill: Aggregating data

Make a table showing the proportion of blacks and whites that are controls and cases.

💡 Expand to see solution

```
prop.table(table(ds$case_control_status, ds$race), margin = 2)
```

	B	W	White
0	0.5396825	0.5156250	0.0000000
1	0.4603175	0.4843750	1.0000000

34.1 Comment:

The `margin` parameter of the `prop.table` command has to be specified in order to get the desired answer: “1 indicates rows, 2 indicates columns.

```
prop.table(table(ds$case_control_status, ds$race), margin = 1)
```

	B	W	White
0	0.67326733	0.32673267	0.00000000
1	0.64444444	0.34444444	0.01111111

```
prop.table(table(ds$case_control_status, ds$race))
```

	B	W	White
0	0.356020942	0.172774869	0.000000000
1	0.303664921	0.162303665	0.005235602

Construct more readable tables with labels using `xtabs`

💡 Expand to see solution

34.2 xtabs table with labels

```
prop.table(xtabs(~case_control_status + race, data = ds), margin = 1)
```

	race		
case_control_status	B	W	White
0	0.67326733	0.32673267	0.00000000
1	0.64444444	0.34444444	0.01111111

Create a count cross table using Tidyverse commands

💡 Expand to see solution

```
ds %>%  
  group_by(case_control_status, race) %>%  
  summarize(n = n()) %>%  
  spread(race, n)
```

`summarise()` has grouped output by 'case_control_status'. You can override using the `.groups` argument.

```
# A tibble: 2 x 4  
# Groups:   case_control_status [2]  
  case_control_status      B      W White  
      <dbl> <int> <int> <int>  
1             0      68     33    NA  
2             1      58     31     1
```

```
addmargins(xtabs(~case_control_status + race, data = ds))
```

	race			
case_control_status	B	W	White	Sum
0	68	33	0	101
1	58	31	1	90
Sum	126	64	1	191

Create a proportion cross table using Tidyverse commands

💡 Expand to see solution

```
ds %>%
  group_by(case_control_status, race) %>%
  summarize(n = n()) %>%
  mutate(prop = n/sum(n)) %>%
  select(-n) %>%
  spread(race, prop)
```

`summarise()` has grouped output by 'case_control_status'. You can override using the `.groups` argument.

```
# A tibble: 2 x 4
# Groups:   case_control_status [2]
  case_control_status      B      W  White
      <dbl> <dbl> <dbl>   <dbl>
1             0 0.673 0.327    NA
2             1 0.644 0.344 0.0111
```

35 Exercise 4

Skill: Summarizing within groups

Apply the `summary` command to the “Gestationalage_sample” within each “Sample_trimester” group.

💡 Expand to see solution

```
f <- split(ds[, "Gestationalage_sample"], ds$Sample_trimester)
sapply(f, summary)
```

```
      1      2      3
Min.  4.934325 16.53800 31.44880
1st Qu. 7.838825 18.45761 35.16305
Median  8.565282 19.72388 37.71093
Mean    8.616799 19.83310 37.37827
3rd Qu.  9.193104 20.69576 39.11360
Max.    13.026958 24.60659 42.09340
```

```
# Or 'tapply' can be used:
tapply(ds$Gestationalage_sample, ds$Sample_trimester, summary)
```

```
$`1`
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 4.934  7.839   8.565   8.617  9.193  13.027
```

```
$`2`
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
16.54  18.46   19.72   19.83  20.70   24.61
```

```
$`3`
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
31.45  35.16   37.71   37.38  39.11   42.09
```

Note: With `split(x, f)`, any missing values in `f` are dropped together with the corresponding values of `x`.

36 Exercise 5: Recoding data

Approach 1

- Implement our dictionaries using look-up tables
 - Use a named vector.

Skill:: Recoding IDs using a dictionary

Create a new subject ID column named “subjectID” where you have used the `DictPer` named vector to recode the original “subject_id” IDs into integer IDs.

```
head(DictPer)
```

```
SUBJ48 SUBJ46 SUBJ28 SUBJ26 SUBJ49 SUBJ19
      40      2      23      38      10      27
```

💡 Expand to see solution

```
a5 <- ds
a5$ID <- DictPer[a5$subject_id]
a5 %>%
  select(subject_id, ID) %>%
  head
```

```
subject_id ID
1      SUBJ48 40
2      SUBJ46  2
3      SUBJ28 23
4      SUBJ26 38
5      SUBJ49 10
6      SUBJ48 40
```

```
head(DictPer)
```

```
SUBJ48 SUBJ46 SUBJ28 SUBJ26 SUBJ49 SUBJ19
      40      2      23      38      10      27
```

37 Recoding data

Approach 2

- Implement our dictionaries using left joins

37.1 Comment

I usually prefer to use a merge command like `left_join` to merge in the new IDs into my data frame.

💡 Expand to see solution

```
key <- data.frame(SubjID = names(DictPer), ID = DictPer)
head(key)
```

```
      SubjID ID
SUBJ48 SUBJ48 40
SUBJ46 SUBJ46  2
SUBJ28 SUBJ28 23
SUBJ26 SUBJ26 38
SUBJ49 SUBJ49 10
SUBJ19 SUBJ19 27
```

```
b5 <- left_join(ds, key, by = c(subject_id = "SubjID"))
b5 %>%
  select(subject_id, ID) %>%
  head()
```

```
subject_id ID
1      SUBJ48 40
2      SUBJ46  2
3      SUBJ28 23
4      SUBJ26 38
5      SUBJ49 10
```


38 Exercise 6

Skill: Filtering rows.

Create a data frame `tri1` containing the records for Trimester 1, and a second data frame `tri2` containing the records for Trimester 2.

 Expand to see solution

```
tri1 <- ds %>%  
  filter(Sample_trimester == 1)  
tri1 %>%  
  select(subject_id, sample_id, Sample_trimester) %>%  
  head()
```

	subject_id	sample_id	Sample_trimester
1	SUBJ48	SAMP149	1
2	SUBJ46	SAMP037	1
3	SUBJ28	SAMP120	1
4	SUBJ48	SAMP149	1
5	SUBJ07	SAMP063	1
6	SUBJ28	SAMP053	1

```
tri2 <- ds %>%  
  filter(Sample_trimester == 2)  
tri2 %>%  
  select(subject_id, sample_id, Sample_trimester) %>%  
  head()
```

	subject_id	sample_id	Sample_trimester
1	SUBJ49	SAMP082	2
2	SUBJ19	SAMP074	2
3	SUBJ10	SAMP121	2
4	SUBJ22	SAMP184	2
5	SUBJ29	SAMP100	2
6	SUBJ19	SAMP074	2

39 Exercise 7

Skill: Selecting columns

Update `tri1` and `tri2` to only contain the three columns “`sample_id`”, “`Sample_trimester`”, “`Gestationalage_sample`”

💡 Expand to see solution

```
tri1 <- tri1 %>%  
  select(sample_id, Sample_trimester, Gestationalage_sample)  
head(tri1)
```

	sample_id	Sample_trimester	Gestationalage_sample
1	SAMP149	1	8.094299
2	SAMP037	1	7.146034
3	SAMP120	1	7.122495
4	SAMP149	1	8.473876
5	SAMP063	1	7.510132
6	SAMP053	1	7.446434

```
tri2 <- tri2 %>%  
  select(sample_id, Sample_trimester, Gestationalage_sample)  
head(tri2)
```

	sample_id	Sample_trimester	Gestationalage_sample
1	SAMP082	2	21.89337
2	SAMP074	2	21.26259
3	SAMP121	2	18.29106
4	SAMP184	2	18.76825
5	SAMP100	2	24.48074
6	SAMP074	2	21.24652

40 Summary

In summary, this book is a work in progress.

References